

---

# Human Computation in Online Video Storytelling

---

Philo van Kemenade

Submitted in partial fulfillment of the requirement of the degree of  
MASTER OF SCIENCE IN COGNITIVE COMPUTING

Goldsmiths College  
University of London

*Supervisor*

Dr. Marian Ursu

Department of Computing  
Goldsmiths, University of London  
New Cross, London SE14 6NW

September 14, 2012

## **Abstract**

Digital video retrieval, filtering and reconfiguration are difficult tasks to solve using current computational techniques. An important cause of this difficulty is the semantic gap between a visual representation and the meaning we address to it. A solution commonly sought in AI research is to reduce the gap by visual analysis and the linking thereof to previously established semantical concepts. These methods often perform poorly on unpredictable content found in the large video libraries of user generated content that make up much of the web's traffic these days. A second way of hunting down meaning in visual content is to step over the gap altogether and ask people directly for a meaningful interpretation one wishes to acquire for an item of content. By accessing many people's interpretations in small bite-sized tasks, collective notions of meaning can be established. This form of accessing human computational power has seen a major increase in attention and application, for a large part because of the increased connectivity of individuals across the web. This thesis explores how tasks involving meaningful interpretation of video content can benefit from the use of Human Computation Systems. In order to test the validity of these approaches 'wePorter' is developed, a system with the purpose of finding local intervals of interest within videos in a larger set of topically related content. We also investigate how such a system can be used for meaningful story reconfiguration. [concluding]



# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
<b>2</b>	<b>The Quest for Meaning in Video</b>	<b>5</b>
2.1	Forms of Meaning in Film and Video . . . . .	5
2.1.1	Meaning in Visuals . . . . .	5
2.1.2	Meaning in Concept . . . . .	5
2.1.3	Meaning in Structure . . . . .	5
2.1.4	Meaning in Annotations . . . . .	5
2.2	Computational Undertakings of the Quest . . . . .	5
2.2.1	Steps towards Meaning: An Overview . . . . .	5
2.2.2	Indexing to enable search . . . . .	5
<b>3</b>	<b>Human Computation towards Visual Meaning</b>	<b>6</b>
3.0.3	The web as platform for creation . . . . .	6
3.1	Deriving meaning from video via Human Factors . . . . .	6
3.2	collaborative filtering . . . . .	7
3.3	A Characterisation of Human Computation Systems . . . . .	7
3.3.1	Purpose . . . . .	7
3.3.2	Motivation . . . . .	7
3.3.3	Task . . . . .	7
<b>4</b>	<b>Interactive Storytelling: From database to data-based</b>	<b>8</b>
4.1	Symbolic approaches . . . . .	8
4.2	Statistic approaches . . . . .	8
4.3	Remixing . . . . .	8
4.3.1	Taking the remix online . . . . .	8
<b>5</b>	<b>Human Computed Stories in wePorter</b>	<b>9</b>
5.1	Introduction . . . . .	9
5.2	User Generated Video Content . . . . .	9
5.3	The Purpose . . . . .	10
5.3.1	Different Goals . . . . .	10
5.4	serving the purpose of unarticulated want . . . . .	11
5.4.1	Storytelling as Structured Recommendation . . . . .	12
5.5	The Motivation . . . . .	12
5.5.1	Information Provision through Online Video . . . . .	13
5.6	The Task . . . . .	14
5.6.1	Between consumer and producer . . . . .	14
5.6.2	Forced Feedback . . . . .	15

5.7	Analysis of the wePorter System . . . . .	15
5.7.1	Landscapes of Interest . . . . .	15
5.7.2	The death of the author . . . . .	15
5.8	Parallel Play . . . . .	15
5.9	Implementation . . . . .	16
<b>6</b>	<b>Evaluation</b>	<b>17</b>
<b>7</b>	<b>Discussion</b>	<b>18</b>
<b>8</b>	<b>Future Directions</b>	<b>19</b>
<b>9</b>	<b>Conclusions</b>	<b>20</b>
	<b>Bibliography</b>	<b>21</b>

## Chapter 1

# Introduction

## Chapter 2

# The Quest for Meaning in Video

### 2.1 Forms of Meaning in Film and Video

#### 2.1.1 Meaning in Visuals

[Video structuring, indexing and retrieval based on global motion wavelet coefficients][1]

#### 2.1.2 Meaning in Concept

#### 2.1.3 Meaning in Structure

[cite bordwell & Thompson: analysis of context dependency]

#### 2.1.4 Meaning in Annotations

[Telop-on-demand: Video structuring and retrieval based on text recognition][9] [Addressing the Challenge of Visual Information Access from Digital Image and Video Libraries][3]

### 2.2 Computational Undertakings of the Quest

#### 2.2.1 Steps towards Meaning: An Overview

Schematized summary of different steps: indexing automatic metadata annotating Human-Driven Labeling Machine-Driven Labeling multimodal feature fusion concept based content based [Relevance feedback: A power tool for interactive content-based image retrieval][14] [The Relative Effectiveness of Concept-based Versus Content-based Video Retrieval][19] [collaborative filtering] [http://en.wikipedia.org/wiki/Collaborative\\_filtering](http://en.wikipedia.org/wiki/Collaborative_filtering)

#### 2.2.2 Indexing to enable search

Because visual data on it's own provides little machine-readable handles to search and find, repositories of multi-media content need to be index to enable search. Within the task of video indexing several approaches are taken

## Chapter 3

# Human Computation towards Visual Meaning

In a recent survey paper, Quinn and Bederson present a taxonomy of Human computation systems[11]. They sketch out the trend of this new method for intelligent problem solving by the increase of academic papers featuring the term ‘human computation’ and its relative ‘crowd-sourcing’. They summarise the myriad of definitions given in recent works by several different authors in two key points:

- “The problems fit the general paradigm of computation, and as such might someday be solvable by computers.”
- “The human participation is directed by the computational system of process”

The first point introduces an interesting question whether storytelling is a computable process. Turing envisioned in his seminal paper computing machinery and human intelligence from 1950 [?] that a computer program would be able to successfully play a game now known as the Turing Test. ...

Quinn and Bederson further present a classification along six dimensions they see as the most salient distinguishing factors:

[table of dimensions?]

### 3.0.3 The web as platform for creation

Many media scholars have written about the role of the web [refs New Media Reader]. Important trend of the web as platform of creation. In terms of video creation for example, the last few years have seen the development of online video editing tools and environments such as popcorn.js, WeVideo and Kaltura.

## 3.1 Deriving meaning from video via Human Factors

[Personalized online document, image and video recommendation via commodity eye-tracking][17]

[VideoReach: an online video recommendation system][10]



## **3.2 collaborative filtering**

The idea of using user's past interactions within a system hosting digital content for the filtering of items that might be of interest is not a new one and usually goes by the name of collaborative filtering. Collaborative filtering can generally take two forms: User-based, Item-based

Information filtering agents and collaborative filtering both attempt to alleviate information overload by identifying which items a user will find worthwhile. I already happening at YouTube

## **3.3 A Characterisation of Human Computation Systems**

### **3.3.1 Purpose**

### **3.3.2 Motivation**

### **3.3.3 Task**

## Chapter 4

# Interactive Storytelling: From database to data-based

### 4.1 Symbolic approaches

[21, 12, 16]

### 4.2 Statistic approaches

### 4.3 Remixing

The reconfiguration of smaller units that carry meaning within themselves is common practice for textual media such as blogs, where it is easy to quote part of another author's writing in a new post [ref].

It needs to be said that some reconfiguration of videos is taking place, but even though it often concerns content that was originally sourced online, much of the creative act of remixing happens offline.

#### 4.3.1 Taking the remix online

examples like: Aaron Koblin (johnny cash project, exquisite corps, etc)

there is even word of a true remix culture.

the availability of multi-media content via the internet has meant a surge in

## Chapter 5

# Human Computed Stories in wePorter

### 5.1 Introduction

This section describes the interaction design of the wePorter system, built to exemplify how human computation can be used in tasks like local video part filtering and semi-automated video reconfiguration. The wePorter system runs an interactive webpage that functions as the main source for data acquisition, presentation of results and general proof of concept. In this chapter the system is analysed along the axes of *Purpose*, *Motivation* and *Task*, that were introduced in the analysis of Human Computational systems in chapter 3. Next to these guidelines for analysis, some remarks are made about the specifics in the functioning of the system. Lastly we discuss the implementation of the web application that is central in wePorter.

### 5.2 User Generated Video Content

[define UGVC]

[http://youtube-global.blogspot.co.uk/2009/05/zoinks-20-hours-of-video-uploaded-every\\_20.html](http://youtube-global.blogspot.co.uk/2009/05/zoinks-20-hours-of-video-uploaded-every_20.html)

<http://youtube-global.blogspot.co.uk/2010/03/oops-pow-surprise24-hours-of-video-all.html>

<http://youtube-global.blogspot.co.uk/2010/11/great-scott-over-35-hours-of-video.html> <http://youtube-global.blogspot.co.uk/2011/05/thanks-youtube-community-for-two-big.html> <http://youtube-global.blogspot.co.uk/2012/01/holy-nyans-60-hours-per-minute-and-4.html>

Since the dawn of YouTube, we've been sharing the hours of video you upload every minute. In 2007 we started at six hours, then in 2010 we were at 24 hours, then 35, then 48, and now...60 hours of video every minute, an increase of more than 25 percent in the last eight months. In other words, you're uploading one hour of video to YouTube every second. Tick, tock, tick, tock that's 4 hours right there!

These astonishing figures of the amount of video that is uploaded to youtube are nothing short of mind blowing, but will most likely sound dated in a matter of years or even months. Looking at the increase of content uploaded to the video platform in

past years, the growth does not seem likely to come to a halt soon [ref table]. All these videos are great for online video junkies, and are increasingly part of the online journalism landscape [13]. At the same time, all these videos being put online beg the question which ones of them to watch.

[table of youtube content uploads]

The increasing amounts of content being put online, lead to an information overload and present serious challenges in search and information retrieval tasks [ref]. There is an increased need for ways of aggregation and filtering. Both of these tasks rely heavily on an at least a shallow understanding of what is presented in these media, which, as we've seen in chapter 2, is a hard problem to solve via current computational techniques. With so much content being uploaded, how can we find our way in the already enormous ocean of online videos?

## 5.3 The Purpose

Searching -¿ IR With more than an hour of new content per second it is no wonder that youtube has come to be viewed as the go-to for online video, much like “the digital video repository for the Internet” that was envisioned by its founders in their first ever blog post [ref <http://youtube-global.blogspot.co.uk/2005/07/greetings-everyone-thanks-for-visiting.html>]. An important activity on video platforms like youtube is searching and much attention has been given to different methods of multimedia search and indexing [refs]. Youtube's acquisition by Google in 2006 underlines the platform's role as a video search engine.

### 5.3.1 Different Goals

Once items have been annotated with tags that reflect the content of a video, these indexes can, along with other meta data of the video, be used for retrieval of videos[ref] in response to textual queries. The effectiveness of such a retrieval task can vary depending on the information that is used in the search algorithm[refs] and the type of content that is searched for [7][more refs]. A third characteristic that determines the effectiveness of a video retrieval system is the goals that users have in their usage of the system. These user goals can vary widely from more to less specific[4]. We expand on this latter point, as it forms an important context for the wePorter system.

#### Direct Navigation

First, a user might be drawn to a video platform by a direct link from an external website. Links might be either in the form of actual hyperlinks or playable embedded videos that are followed through to the platform. Navigations via such links form a direct mapping between a user's intention to the desired piece of content. In this case, users have a very specific reason to come and watch. Their desire, at least of knowing the contents of the video, is satisfied after the viewing. YouTube's system engineers call this way of video viewing *direct navigation*[4].

#### Search and Goal-oriented browse

When users have not obtained a direct link to a potentially relevant piece of content, they might still have a specific goal in mind when visiting an online video platform. Reasons to visit might be the wish to see a particular music video or to find an instance of a series

by a particular producer. This goal of discovering a rather specific video is referred to as *search and goal-oriented browse*. Provided that the desired piece of content exists and the video platform has an appropriate search function in place, these ‘narrow queries’, will result in a result set of search results from which the user is likely to handpick the sought-after result fairly quickly. Here the user’s desired result often lies within a single item of content. Perhaps a few hit and misses are required, but after a couple of clicks the required video is found.

### **Unarticulated Want**

A different scenario emerges altogether when a user approaches a video platform with a more broad and open ended motive. Think for example of someone who wants to know what happened at a large music festival she recently attended, or someone who couldn’t make it to a large demonstration and would like to get a sense of the atmosphere. This kind of ‘broad queries’ returns a result set of content in which a users will probably consider many items as a successful retrieval. Furthermore, one could even say that the desired result of a user’s query is spread across the multiple pieces of content. By traversing the space of different videos in the result set, users interactively construct the answers to their own queries.

This kind of navigation through a space of related content is common practice on large spaces of linked data[ref]. It has been found that YouTube’s related video recommendation functionality, which recommends videos that are related to the video currently being watched, is one of the most important view sources of videos. In fact, traffic received from these recommendations is the main source of views for the majority of videos on YouTube [20]. Features derived from users’ navigations such as ‘click-through rate’ have been used to improve content-based video recommendation [18].

Goal of a person’s query in this kind of navigation is no longer defined in a single returnable item of content or even a containable set of items. Rather, the interactive pathway through the a set of interesting bits of content is what represents a user’s aim. This broader exploratory goal of finding different parts of relevant content has been termed ‘unarticulated want’[4].

Interactivity is generally agreed to play an important role in the task of video retrieval, as is reflected by the separate category in the annual TRECvid challenge for interactive video retrieval[15]. Several works have indicated the importance of interactivity in the task of video retrieval to filter through a set of initially returned results [5, 2, 5, 6]. As most of these systems are aimed at retrieval of clearly specified intentions, for example in the TRECvid retrieval task, the need for interactive exploration is even more apparent for the broader oriented goal of users engaged in unarticulated want.

## **5.4 serving the purpose of unarticulated want**

The answer to a user’s query now lies as much in the journey through the content as in the returned content itself. This self-constructed story is an important concept that wePorter capitalises on, as will soon become apparent.

The task at hand of recommending a larger group of interesting videos is radically different compared to the more narrow queries that could be answered by a small set of true positives in an information retrieval task. Besides the spread of the searched for result across different pieces of content, there is a second important difference that lies in the nature of pieces of UGVC.

Users with broad expectations will not only want to be presented with a set of relevant items from a complete repository, they are also looking for the most interesting parts within these relevant items. This issue is particular to time-based media, and especially relevant for video. Other temporal media, like audio in general and music in particular, have less of a need for segmentation because of their common usage in multi-media applications. People usually tend to listen to a song entirely and if they wish to experience an album in part, constituent songs are already units on their own that can easily be reconfigured.

Because of the raw, unedited nature of the majority of UGVC [ref yt and the news] it is desirable to establish local recommendations that point to parts within a video content that are of particular interest. Whereas digital music albums shared online consist of a collection of songs that can each easily be made to stand alone, video currently suffers from a less malleable identity online. Online videos are currently much like black boxes that can be played, paused, rated, commented on, tagged and shared only in its entirety. What if a piece of raw, unedited UGVC features something spectacular for ten seconds halfway along its timeline, but shows much of the same for the rest of the time? Answering this question will be the first part of the purpose of the wePorter system.

#### 5.4.1 Storytelling as Structured Recommendation

explanation why it is also good to structure the filtered pieces of content into a new configuration. [TODO]

The ten significant seconds in a two-minute video become a needle in a haystack when an initial set of videos relating to your query includes tens to hundreds of possibly relevant videos with lengths between some tens of seconds and a couple of minutes. The aggregation and reconfiguration of several of these ‘needles’ into a meaningful new whole is another non-trivial task. We present wePorter as a test case for new methods that address both these issues of information overload in video libraries of UGVC. More precisely, wePorter’s purpose is two-folded:

From a set of topically related unedited user-generated videos:

1. Find localised intervals of interest within each of the source videos
2. Find a meaningful structure for the reconfiguration of interesting video parts

Idea 1: towards interest based filtering via meaning

Idea 2: skip the semantic gap and directly model human attentional behaviour

many people watching: Clicking from one video to the next (choosing from a set of related videos) these inter-video links could be seen as indicators for relatedness and relevance, much like google’s page rank algorithm use links across webpages to establish a notion of the most significant site on a particular topic.

There is an important difference here though. Whereas the links used by Google’s search algorithms are embedded in machine readable hyperlinks, the path of clicking on from one video to the next is a characteristic of a person’s interaction.

differences: public, readable // private, non readable conscious choice // unconscious result of interaction Concluding can be consciously put in place by several people at large scale // dependent on real ‘human’ traffic.

## 5.5 The Motivation

How to get a group of unrelated people to contribute their efforts to solving the tasks set in our two-folded purpose? This section looks at the reasons people might have to contribute

their computational powers to a system with a purpose like wePorter. Looking at the way people engage with online video content on platforms like youtube, we identify patterns in their behaviour that can be matched to a task in a human computation system. This behaviour that is characterised by a more active role in multimedia consumption, can be seen as a larger trend in the development of new media. The end of this section indicates how the motivations of users of the wePorter system can link in with this larger trend.

### 5.5.1 Information Provision through Online Video

Since the proliferation of mobile video recording devices, it has become common practice for large-scale (semi-)public events to be covered in UGVC that gets uploaded to the web. While some are critical[8] about the often heralded democratization and empowerment of people by the new media production and distribution tools, it is clear that the UGVC at places like youtube attracts a lot of traffic from people looking to be informed about recent events. Afterall UGVC can have its advantages over traditional video news production, especially for unexpected events where traditional media production does not have the immediacy of user-generated ‘reports’ recorded by coincidental passersby.

In a recent study as part of the the Pew Research Centers Project for Excellence in Journalism, for a period of 15 months the most popular video’s from YouTube’s ‘News and Politics’ category are analyzed[13]. The authors of the study exemplify the power UGVC can have in news provision by showcasing frequently viewed videos detailing scenes from the earthquake and subsequent tsunami that hit Japan in March 2011. The week following the disaster, the 20 most viewed news-related videos on YouTube were all focused on the catastrophic event and were viewed more than 96 million times. Most of these videos were recorded by individuals who happened to be in the affected areas when the disaster struck, either uploaded by themselves, or by TV channels who appropriated the content. The study furthermore reports that in the investigated period, the most searched term of the month was a news-related event 5 out of 15 months.

While the journalism study above focusses on videos with the ‘News & Politics’ label, information provision about current event might span a larger set of categories. Someone looking for footage in order to get a sense of the atmosphere at a recent music festival or public demonstration, might very well find relevant videos in categories like ‘Entertainment’, ‘Travel & Events’ or ‘Nonprofits & Activism’. Across all of these categories, we are able to find examples of UGVC, uploaded in the period following up newsworthy events.

The wePorter system focusses on these kinds of topically related sets that people are currently exploring interactively by browsing from one video to the next. This way of navigation is an intermediary between the goals of *goal-oriented browse* and *unarticulated want*. The unarticulated want is now encapsulated by the event but users still roam freely within this topicalized set of content. By navigating from video to video, watching some and skipping others, users leave attentional traces that give valuable insight into a user’s intentional standpoint.

It is this kind of interactions that are already taking place at a large scale that we like to make use of in the wePorter system. Motivated by the wish to explore informative content, users of system appropriately appropriately set up, will instinctively and implicitly contribute their human knowledge. ... and a task is set up such that their interactions are recorded as contributions to solving small instances of a larger problem we have two of the three ingredients of our human computation system in place.

## 5.6 The Task

This would fall into what Quinn and Bederson would label as ‘implicit work’ [11].

### 5.6.1 Between consumer and producer

Intro We start by explaining the motivations out of which the design was born, followed by a detailed explanation of the interface and analysis of some of the design choices that were made in its development.

Between Couch Potato and Video Editor The interaction here presented is the result of a design process motivated by the wish to solve a challenge:

How to get massive amounts of people collaboratively contributing to the shape of a video story?

An answer to this question should balance extremes in multiple dimensions: creation vs consumption active vs passive advanced control vs easy access

A central idea in the wePorter system is that video consumption and video creation might be more closely related than traditionally thought. In fact, relating these two aspects is a key element in the interaction design here presented.

Traditionally, models of media productions, distribution and consumption have usually assumed a strict division between the roles of producers and consumers. [need for New Media studies ref] The classical image of television viewers is that of passive couch potatoes, comfortably leaned back on the sofa in their living room. Literally miles apart from them are content producers like reporters, presenters, cameramen, directors, editors and so on, who capture, form and present the content that is viewed on screen.

Thanks to the creative use of telecommunications technology, some television programmes have embraced more interactivity between the two roles. Formats like the numerous talent shows that let people vote by text or gameshows where people play by calling in by telephone, are examples where consumers influence what happens onscreen, but roles of creator and consumer stay effectively unchanged. Creators are still responsible for the production of content shown onscreen. In living up to this responsibility, they

Video Editor Couch Potato creator consumer active passive making viewing active passive difficult easy

Parallel player

In the web-based interface two video players are presented at once. When a user initiates an interaction session by hitting play, the two videos start playing. Both videos are now visible to the user, but only one of the audio tracks is played. Users control which player’s audio track is being played by moving the mouse cursor over the player they wish to hear. Apart from triggering the audio of a video this way, mouse-over movement also brings focus to a video by displaying it clear, while video out of focus is silent and slightly dimmed.

For every 10th of a second, we record which video a user attends to. Note that we never explicitly ask anyone to point at the video that is most interesting. Users are simply instructed as to how the interface works and then left to explore the video as they like. By recording users’ behaviour this way, we achieve a detailed insight as to which of a pair of videos a user has attended to at what time.

Positioning two video’s one on top of the other, might inflict a bias for users in their attentional behaviour. It might be the case videos on the top are systematically more attended to than videos displayed below. We’ve experimented to see whether such positioning bias effects occur and report on this in section ?? [\*TODO ref].

Sequences of video parts



The video sequences played by the two players consist of distinct parts of videos from a topic-centred video database. This means all content that is aggregated in the interaction is focussed around a single topic, in this example the event of the Queen of England's Diamond Jubilee celebrations in London. The sequences have some other important characteristics: Sequences are formed out of the same number of video parts (in this example 6) Video parts in sequences are all of the same duration (in this example 10 seconds) It follows that all sequences have equal duration (in this example 1 minute).

### 5.6.2 Forced Feedback

By making data acquisition from interaction an intrinsic part of a human computation system [\*need umbrella term], submitting data to the system will become second nature to users and can even take place without the users being aware of it. -¿ curation through interaction

## 5.7 Analysis of the wePorter System

### 5.7.1 Landscapes of Interest

The distinction to be made, between interesting intervals on one hand and less striking parts of a video on the other is not likely to be a very strict one. Afterall, an unedited video captures a single stretch of space and time, so any event that is of particular interest will unlikely have hard cut-off points in time. Rather, if we see interest as a function of time in a particular video, we would expect a somewhat continuous flowing line with spikes every now and then when an interesting event occurs.

Looking at interest at a more global level, aggregating over a large group of users, would perhaps even a more smooth landscape of interest. This kind of data could reveal mountains and valleys that can be used for interest-based segmentation. From the thus segmented parts, the ones with high interest score can be returned as the salient parts within a video.

### 5.7.2 The death of the author

ref: Barthes - Image music Sound here: a lack of clear narrator of the written story

in wePorter: the lack of a centralised creator, authoring a work consciously and deliberately.

This new form of creation could at first sight be seen as non-authorship, but is rather a collective authorship.

Several recent developments now make possible this collective form of authorship: - the proliferation of tools that enable individual authorship of multimedia content - the increased connectivity of these devices, which gives them the capacity to make the created media accessible to others - the platforms for hosting multimedia (MM) content. This point is strongly related to the previous one. In their development they form a chicken and egg relation. Initially it's hard to imagine one without the other, but once matured, they [influence each other positively]. - the Methods to link, mix, aggregate and modify content online.

## 5.8 Parallel Play

A new method for preference elicitation in time-based multi-media content.

## 5.9 Implementation

argmin  
 $x$

---

**Algorithm 1** My algorithm

---

```
1: procedure EUCLID( $a, b$ )                                ▷ The g.c.d. of a and b
2:    $r \leftarrow a \bmod b$ 
3:   while  $r \neq 0$  do                                    ▷ We have the answer if r is 0
4:      $a \leftarrow b$ 
5:      $b \leftarrow r$ 
6:      $r \leftarrow a \bmod b$ 
7:   end while
8:   return  $b$                                              ▷ The gcd is b
9: end procedure
```

---

---

**Algorithm 2** Generate Sequences Random Shuffled

---

```
1: procedure FILL SEQUENCES                                ▷ The g.c.d. of a and b
2:    $r \leftarrow a \bmod b$ 
3:   while  $r \neq 0$  do                                    ▷ We have the answer if r is 0
4:      $a \leftarrow b$ 
5:      $b \leftarrow r$ 
6:      $r \leftarrow a \bmod b$ 
7:   end while
8:   return  $b$                                              ▷ The gcd is b
9: end procedure
```

---

## Chapter 6

# Evaluation

This is the Evaluation section.

## Chapter 7

# Discussion

## Chapter 8

# Future Directions

This is the Future Directions chapter.

## Chapter 9

# Conclusions

This is the conclusions chapter.

# Bibliography

- [1] E. Bruno and D. Pellerin. Video structuring, indexing and retrieval based on global motion wavelet coefficients. *Pattern Recognition, 2002. Proceedings. 16th International Conference on*, 3:287–290 vol. 3, 2002.
- [2] M. Christel and N. Moraveji. Finding the right shots: assessing usability and performance of a digital video library interface. *Proceedings of the 12th annual ACM international conference on Multimedia*, pages 732–739, 2004.
- [3] M.G. Christel and R.M. Conescu. Addressing the challenge of visual information access from digital image and video libraries. *Proceedings of the 5th ACM/IEEE-CS joint conference on Digital libraries*, pages 69–78, 2005.
- [4] J. Davidson, B. Liebald, J. Liu, P. Nandy, T. Van Vleet, U. Gargi, S. Gupta, Y. He, M. Lambert, and B. Livingston. The YouTube video recommendation system. *Proceedings of the fourth ACM conference on Recommender systems*, pages 293–296, 2010.
- [5] O. De Rooij, C G M Snoek, and M Worring. Query on demand video browsing. *Proceedings of the 15th international conference on Multimedia*, pages 811–814, 2007.
- [6] O. De Rooij, C G M Snoek, and M Worring. Balancing thread based navigation for targeted video search. *Proceedings of the 2008 international conference on Content-based image and video retrieval*, pages 485–494, 2008.
- [7] L Hollink, G P Nguyen, D C Koelma, A Th Schreiber, and M Worring. Assessing user behaviour in news video retrieval. *IEE Proceedings - Vision, Image, and Signal Processing*, 152(6):911, 2005.
- [8] Anna Maria Jönsson and Henrik Örnebring. USER-GENERATED CONTENT AND THE NEWS. *Journalism Practice*, 5(2):127–144, April 2011.
- [9] H. Kuwano, Y. Taniguchi, H. Arai, M. Mori, S. Kurakake, and H. Kojima. Telop-on-demand: Video structuring and retrieval based on text recognition. *Multimedia and Expo, 2000. ICME 2000. 2000 IEEE International Conference on*, 2:759–762 vol. 2, 2000.
- [10] T. Mei, B. Yang, X.S. Hua, L. Yang, S.Q. Yang, and S. Li. VideoReach: an online video recommendation system. *Proceedings of the 30th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 767–768, 2007.
- [11] A.J. Quinn and B.B. Bederson. Human computation: a survey and taxonomy of a growing field. *Proceedings of the 2011 annual conference on Human factors in computing systems*, pages 1403–1412, 2011.

- [12] Pablo Cesar, Dick C. A. Bulterman, Vilmos Zsombori, Ian Kegel, Rodrigo Laiola Guimaraes. Creating Personalized Memories from Social Events: Community-Based Support for Multi-Camera Recordings of School Concerts. Technical report, August 2011.
- [13] Tom Rosenstiel and Amy Mitchell. YouTube & the News. Technical report, July 2012.
- [14] Y. Rui, T.S. Huang, M. Ortega, and S. Mehrotra. Relevance feedback: A power tool for interactive content-based image retrieval. *Circuits and Systems for Video Technology, IEEE Transactions on*, 8(5):644–655, 1998.
- [15] A.F. Smeaton, P. Over, and W. Kraaij. Evaluation campaigns and TRECVID. *Proceedings of the 8th ACM international workshop on Multimedia information retrieval*, pages 321–330, 2006.
- [16] Marian F Ursu, Vilmos Zsombori, John Wyver, Lucie Conrad, Ian Kegel, and Doug Williams. Interactive documentaries. *Computers in Entertainment*, 7(3):1, September 2009.
- [17] S. Xu, H. Jiang, and F. Lau. Personalized online document, image and video recommendation via commodity eye-tracking. *Proceedings of the 2008 ACM conference on Recommender systems*, pages 83–90, 2008.
- [18] B. Yang, T. Mei, X.S. Hua, L. Yang, S.Q. Yang, and M. Li. Online video recommendation based on multimodal fusion and relevance feedback. *Proceedings of the 6th ACM international conference on Image and video retrieval*, pages 73–80, 2007.
- [19] M. Yang, B.M. Wildemuth, and G. Marchionini. The relative effectiveness of concept-based versus content-based video retrieval. *Proceedings of the 12th annual ACM international conference on Multimedia*, pages 368–371, 2004.
- [20] R. Zhou, S. Khemmarat, and L. Gao. The impact of YouTube recommendation system on video views. *Proceedings of the 10th annual conference on Internet measurement*, pages 404–410, 2010.
- [21] V. Zsombori, M. Frantzis, R.L. Guimaraes, M.F. Ursu, P. Cesar, I. Kegel, R. Craigie, and D. Bulterman. Automatic generation of video narratives from shared UGC. *Proceedings of the ACM Conference on Hypertext and Hypermedia*, pages 325–334, 2011.